

Enterprise Document RAG Sample

This document simulates a telecom enterprise knowledge base.

The system pipeline includes document parsing, text cleaning, chunk splitting, BGE embedding, hybrid retrieval, BGE reranking, refusal control, and DeepSeek answer generation.

BM25 is useful for exact terms such as SLA, network topology, fault diagnosis, and RAGAS.

Vector retrieval is useful when the user question is semantically similar but wording is different.

Reranker is used after initial retrieval to move the most useful evidence chunks to the top.

When the retrieved evidence is not reliable, the system should refuse to answer.